

Chenhao Tan

Associate Professor
University of Chicago
Department of Computer Science
5460 University Avenue
Chicago, IL 60637

Email: chenhao@chenhaot.com
Homepage: <https://chenhaot.com>

Research Interests

Human-centered AI, AI & scientific discovery, AI safety, AI & democratic processes.

Professional Experience

Associate Professor, 2025-present
University of Chicago, Chicago, IL

Assistant Professor, 2021-2024
University of Chicago, Chicago, IL

Assistant Professor, 2017-2020
University of Colorado, Boulder, CO

Research Associate, 2016-2017
University of Washington, Seattle, WA
Supervisor: Noah Smith.

Education

Ph.D. in Computer Science, 2010-2016
Cornell University, Ithaca, NY
Advisor: Lillian Lee.

B.E. in Computer Science, 2006-2010
Tsinghua University, Beijing, China
Advisor: Jie Tang.

B.E. in Economics (second major), 2007-2010
Tsinghua University, Beijing, China

Selected Publications

Qirun Dai, Xiao Liu, Jiawei Zhang, Dylan Zhang, Hao Peng, and *Chenhao Tan*. Towards a Universal Causal Reasoner. NeurIPS 2026 (Spotlight).

Yuyang Jiang, Chacha Chen, Teng Wu, Liwen Sun, Han Liu, Shi Feng, and *Chenhao Tan*. Collaborative Disagreement Resolution for Scalable Oversight. ICML 2026.

Xiaoyan Bai, Alexander Baumgartner, Haojia Sun, Ari Holtzman, and *Chenhao Tan*. The Story is Not the Science: Execution-Grounded Evaluation of Mechanistic Interpretability Research. AI4Science Workshop at ICML 2026 (**Best Poster Award**).

Aniket Vashishtha, Qirun Dai, Hongyuan Mei, Amit Sharma, *Chenhao Tan*, and Hao Peng. Executable Counterfactuals: Improving LLMs' Causal Reasoning Through Code. ICLR 2026.

Haokun Liu*, Yangqiaoyu Zhou*, Mingxuan Li*, Chenfei Yuan, and *Chenhao Tan*. Literature Meets Data: A Synergistic Approach to Hypothesis Generation. ACL 2025.

Ari Holtzman and *Chenhao Tan*. Prompting as Scientific Inquiry. NeurIPS 2025 (Position Papers Track).

Xiaoyan Bai, Ike Peng, Aditya Singh, and *Chenhao Tan*. Concept Incongruence: An Exploration of Time and Death in Role Playing. NeurIPS 2025.

Chao-Chun Hsu, Ziad Obermeyer, and *Chenhao Tan*. A machine learning model using clinical notes to identify physician fatigue. Nature Communications 2025.

Yangqiaoyu Zhou, Haokun Liu, Tejes Srivastava, Hongyuan Mei, and *Chenhao Tan*. Hypothesis Generation with Large Language Models. NLP4Science Workshop at EMNLP 2024.

Karen Zhou, Alexander A. Meitus, Milo Chase, Grace Wang, Anne Mykland, William Howell, and *Chenhao Tan*. Quantifying the Uniqueness of Donald Trump in Presidential Discourse. PNAS Nexus 2024.

Chacha Chen, Shi Feng, Amit Sharma, and *Chenhao Tan*. *Machine Explanations and Human Understanding*. Transactions on Machine Learning Research (TMLR); FAccT 2023; **Best Paper Award** at the ICML 2022 Workshop on Human-Machine Collaboration and Teaming.

Han Liu, Yizhou Tian, Chacha Chen, Shi Feng, Yuxin Chen, and *Chenhao Tan*. *Learning Human-Compatible Representations for Case-Based Decision Support*. ICLR 2023.

Mourad Heddaya, Solomon Dworkin, *Chenhao Tan*, Rob Voigt, and Alexander Zentefis. *Language of Bargaining*. ACL 2023.

Yangqiaoyu Zhou, Yiming Zhang, and *Chenhao Tan*. *FLamE: Few-shot Learning from Natural Language Explanations*. ACL 2023.

Vivian Lai, Chacha Chen, Alison Smith-Renner, Q. Vera Liao, and *Chenhao Tan*. *Towards a Science of Human-AI Decision Making: A Survey of Empirical Studies*. FAccT 2023.

Yiming Zhang, Shi Feng, and *Chenhao Tan*. *Active Example Selection for In-Context Learning*. EMNLP 2022.

Chao-Chun Hsu and *Chenhao Tan*. *Decision-focused Summarization*. EMNLP 2021.

Ramaravind Kommiya Mothilal, Amit Sharma, and *Chenhao Tan*. *Explaining Machine Learning Classifiers through Diverse Counterfactual Explanations*. FAccT 2020.

Brian Lubars and *Chenhao Tan*. *Ask not what AI can do for you, but what AI should do: Towards a framework of task delegability*. NeurIPS 2019 (Spotlight).

Vivian Lai and *Chenhao Tan*. *On Human Predictions with Explanations and Predictions of Machine Learning Models: A Case Study on Deception Detection*. FAccT 2019.

Chenhao Tan, Vlad Niculae, Cristian Danescu-Niculescu-Mizil, and Lillian Lee. *Winning arguments: Interaction dynamics and persuasion strategies in good-faith online discussions*. WWW 2016.

Chenhao Tan and Lillian Lee. *All Who Wander: On the Prevalence and Characteristics of Multi-community Engagement*. WWW 2015.

Selected Awards and Grants

Best Poster Award at the ICML AI4Science Workshop (2026)
Outstanding Senior Area Chair at ACL (2023)
Sloan Research Fellowship (2023)
Best Paper Award at the ICML Workshop on Human-Machine Collaboration and Teaming (2022)
Google and J.P. Morgan Research Awards (2022)
NSF CAREER Award (2020)
IBM, Amazon, and Salesforce Research Awards (2020)
Best Paper Award at the ACL RepL4NLP Workshop (2018)
Facebook Fellowship (2015)

Selected Service

Co-organizer of the AI and Social Science Conference at the University of Chicago (2026)
Steering Committee at FAccT (2023-2026)
General Chair at the ACM Conference on Fairness, Accountability, and Transparency (FAccT 2023)
Program Chair at the International Conference on Computational Social Science (IC2S2 2022)

Full List of Awards and Grants

Best Poster Award at the ICML AI4Science Workshop (2026)
AI Research Area Award: New Forms of Socio-Cognitive AI, University of Chicago (2025)
AI Research Area Award: AI and the Workplace, University of Chicago (2025)
Sloan Research Fellowship (2023)
Outstanding Senior Area Chair at ACL (2023)
NSF-CSIRO: HCC: Small: From Legislations to Action: Responsible AI for Climate Change (2023)
Best Paper Award at the ICML Workshop on Human-Machine Collaboration and Teaming (2022)
Google Research Scholar Award (2022)
J.P. Morgan Faculty Research Award (2022)
Data and Democracy Initial Grant: Partisan Polarization and Political Dysfunction (co-PI, 2022)
Data and Democracy Initial Grant: Understanding and Auditing Content Moderation Policies (co-PI, 2022)
CDAC Discovery Challenge: AI-Driven Tutorials for Radiology Student Training (lead PI, 2021)

FAI: Towards Adaptive and Interactive Post hoc Explanations, National Science Foundation (lead PI, 2021)

NSF CAREER Award (2020)

IBM Faculty Award (2020)

Amazon Research Award (2020)

Salesforce Research Award (2020)

AI-Society EAGER: Explanations through Diverse, Feasible, and Interactive Counterfactuals, National Science Foundation (sole PI, 2019)

CRII: Harnessing Machine Learning to Improve Human Decision Making: A Case Study on Deceptive Detection, National Science Foundation (sole PI, 2019)

CHS Small: A Genealogical Framework to Understand the Emergence of Online Groups, National Science Foundation (co-PI, 2019)

Best Paper Award at the ACL RepL4NLP Workshop (2018)

BIGDATA Large: Automating Analysis and Feedback to Improve Mathematics Teachers' Classroom Discourse, National Science Foundation (co-PI, 2018)

Facebook Global Literacy Challenges, final list (2018)

Facebook Research Gift (2018)

NVIDIA GPU Grant (2017)

Postdoc Research Award, University of Washington (2017)

Selected to participate in Heidelberg Laureate Forum (2016)

OARU Travel Award for Heidelberg Laureate Forum (2016)

Facebook Fellowship (2015)

Telluride Residential Scholarship (2012-2016)

Yahoo! Key Scientific Challenges (2012)

KDD Student Travel Award (2011, 2010)

Best Undergraduate Thesis of Tsinghua University (Top 1 among ~160) (2010)

Excellent Undergraduate of Tsinghua University (Top 58 among 3144) (2010)

Excellent Undergraduate of Beijing (2010)

National Scholarship (Top 3 among ~160) (2009, 2008)

Full List of Publications

[Google scholar profile]

Qirun Dai, Xiao Liu, Jiawei Zhang, Dylan Zhang, Hao Peng, and *Chenhao Tan*. Towards a Universal Causal Reasoner. NeurIPS 2026 (Spotlight).

Nathan Delisle and *Chenhao Tan*. Bypassing PC₁ Makes SAEs More Reproducible. NeurIPS 2026.

Zekai Tong, Ruiyao Xu, Aryan Shrivastava, *Chenhao Tan*, and Ari Holtzman. The Text Uncanny Valley: Non-Monotonic Performance Degradation in LLM Information Retrieval. NeurIPS 2026 (Evaluation and Datasets Track).

Xiao Liu, Junsol Kim, Shiyang Lai, Jacy Reese Anthis, Siyang Wu, *Chenhao Tan*, and James Evans. Introspection Tools Help LLMs Understand and Control Themselves. NeurIPS 2026 (Position Papers Track).

Sanjana Ramprasad, Pranav Mani, *Chenhao Tan*, Byron C. Wallace, and Elisa Ferracane. Many Benign Errors are Better than a Few Severe Ones: Evaluating Hallucination Severity. NeurIPS 2026 (Evaluation and Datasets Track).

Mingxuan Li, Qirun Dai, Heran Wang, and *Chenhao Tan*. Emergent Misalignment Is Not Magical. EMNLP 2026.

Andrew Ren, Haokun Liu, and *Chenhao Tan*. Robust Failure, Conservative Repair: Textual Knowledge Distillation from Cross-Model Failures. EMNLP 2026.

Xiaoyan Bai, Alexander Baumgartner, Haojia Sun, Ari Holtzman, and *Chenhao Tan*. The Story is Not the Science: Execution-Grounded Evaluation of Mechanistic Interpretability Research. AI4Science Workshop at ICML 2026 (**Best Poster Award**).

Yuyang Jiang, Chacha Chen, Teng Wu, Liwen Sun, Han Liu, Shi Feng, and *Chenhao Tan*. Collaborative Disagreement Resolution for Scalable Oversight. ICML 2026.

Karen Zhou and *Chenhao Tan*. AutoChecklist: Composable Pipelines for Checklist Generation and Scoring with LLM-as-a-Judge. ACL 2026 (Demo Track).

Mingxuan Li, Hanchen Li, and *Chenhao Tan*. HypoEval: Hypothesis-Guided Evaluation for Natural Language Generation. ACL 2026.

Cristina Garbacea, Heran Wang, and *Chenhao Tan*. Personalized Benchmarking: Evaluating LLMs by Individual Preferences. Findings of ACL 2026.

Mourad Heddaya, Manley Roberts, Rohan Wadhawan, and *Chenhao Tan*. When Internalization Fails: Finding Better Targets for Reasoning Compression. Findings of ACL 2026.

Lan Gao, Abani Ahmed, Oscar Chen, Margaux Reyl, Zayna Cheema, Nick Feamster, *Chenhao Tan*, Kurt Thomas, and Marshini Chetty. Governance of AI-Generated Content: A Case Study on Social Media Platforms. CHI 2026.

Aniket Vashishtha, Qirun Dai, Hongyuan Mei, Amit Sharma, *Chenhao Tan*, and Hao Peng. Executable Counterfactuals: Improving LLMs' Causal Reasoning Through Code. ICLR 2026.

Xiao Liu, Haokun Liu, and *Chenhao Tan*. Heuristic-Based Ideation for Guiding LLMs Toward Structured Creativity. ICLR 2026 (Blog Track).

Ari Holtzman and *Chenhao Tan*. Prompting as Scientific Inquiry. NeurIPS 2025 (Position Papers Track).

Xiaoyan Bai, Ike Peng, Aditya Singh, and *Chenhao Tan*. Concept Incongruence: An Exploration of Time and Death in Role Playing. NeurIPS 2025.

Harvey Yiyun Fu, Aryan Shrivastava, Jared Moore, Peter West, *Chenhao Tan*, and Ari Holtzman. AbsenceBench: Language Models Can’t Tell What’s Missing. NeurIPS 2025 (Datasets and Benchmarks Track, Spotlight).

Ziyu Chen, Junfei Sun, Chenxi Li, Tuan Dung Nguyen, Jing Yao, Xiaoyuan Yi, Xing Xie, *Chenhao Tan*, and Lexing Xie. MoVa: Resources for Generalizable Classification of Human Morals and Values. EMNLP 2025.

Karen Zhou, John Giorgi, Pranav Mani, Peng Xu, Davis Liang, and *Chenhao Tan*. From Feedback to Checklists: Grounded Evaluation of AI-Generated Clinical Notes. EMNLP 2025 (Industry Track).

Yuyang Jiang, Chacha Chen, Shengyuan Wang, Feng Li, Zecong Tang, Benjamin M. Mervak, Lydia Chelala, Christopher M. Straus, Reve Chahine, Samuel G. Armato III, and *Chenhao Tan*. CLEAR: A Clinically-Grounded Tabular Framework for Radiology Report Evaluation. Findings of EMNLP 2025.

Chao-Chun Hsu, Ziad Obermeyer, and *Chenhao Tan*. A machine learning model using clinical notes to identify physician fatigue. Nature Communications 2025.

Dang Nguyen and *Chenhao Tan*. On the Effectiveness and Generalization of Race Representations for Debiasing High-Stakes Decisions. COLM 2025.

Haokun Liu*, Yangqiaoyu Zhou*, Mingxuan Li*, Chenfei Yuan, and *Chenhao Tan*. Literature Meets Data: A Synergistic Approach to Hypothesis Generation. ACL 2025.

Jacy Anthis, Kristian Lum, Michael Ekstrand, Avi Feller, and *Chenhao Tan*. The Impossibility of Fair LLMs. ACL 2025.

Lan Gao, Oscar Chen, Rachel Lee, Nick Feamster, *Chenhao Tan*, and Marshini Chetty. “I Cannot Write This Because It Violates Our Content Policy”: Understanding Content Moderation Policies and User Experiences in Generative AI Products. USENIX Security 2025.

Chacha Chen, Han Liu, Jiamin Yang, Benjamin M. Mervak, Bora Kalaycioglu, Grace Lee, Emre Cakmakli, Matteo Bonatti, Sridhar Pudu, Osman Kahraman, Gul Gizem Pamuk, Aytekin Oto, Aritrick Chatterjee, and *Chenhao Tan*. Can Domain Experts Rely on AI Appropriately? A Case Study on AI-Assisted Prostate Cancer MRI Diagnosis. FAccT 2025.

Mourad Heddaya, Kyle MacMillan, Anup Malani, Hongyuan Mei, and *Chenhao Tan*. CaseSumm: A Large-Scale Dataset for Long-Context Summarization from U.S. Supreme Court Opinions. Findings of NAACL 2025.

Yuyang Jiang, Chacha Chen, Dang Nguyen, Benjamin M. Mervak, and *Chenhao Tan*. GPT-4V Cannot Generate Radiology Reports Yet. Findings of NAACL 2025.

Yangqiaoyu Zhou, Haokun Liu, Tejes Srivastava, Hongyuan Mei, and *Chenhao Tan*. Hypothesis Generation with Large Language Models. NLP4Science Workshop at EMNLP 2024.

Mourad Heddaya, Qingcheng Zeng, *Chenhao Tan*, Rob Voigt, and Alexander Zentefis. Causal Micro-Narratives. Workshop on Narrative Understanding at EMNLP 2024.

Emre Kiciman, Robert Ness, Amit Sharma, and *Chenhao Tan*. Causal Reasoning and Large Language Models: Opening a New Frontier for Causality. Transactions on Machine Learning Research (TMLR) 2024.

Karen Zhou, Alexander A. Meitus, Milo Chase, Grace Wang, Anne Mykland, William Howell, and Chenhao Tan. Quantifying the Uniqueness of Donald Trump in Presidential Discourse. PNAS Nexus 2024.

Tianyu Cao, Natraj Raman, Danial Dervovic, and Chenhao Tan. Characterizing Multimodal Long-form Summarization: A Case Study on Financial Reports. COLM 2024.

Chao-Chun Hsu, Erin Bransom, Jenna Sparks, Bailey Kuehl, Chenhao Tan, David Wadden, Lucy Lu Wang, and Aakanksha Naik. CHIME: LLM-Assisted Hierarchical Organization of Scientific Studies for Literature Review Support. Findings of ACL 2024.

Zining Zhu, Hanjie Chen, Xi Ye, Chenhao Tan, Ana Marasović, Sarah Wiegrefe, and Veronica Qing Lyu. Explanation in the Era of Large Language Models. NAACL 2024 (Tutorial).

Brennan Schaffner, Arjun Nitin Bhagoji, Siyuan Cheng, Jacqueline Mei, Jay L. Shen, Grace Wang, Marshini Chetty, Nick Feamster, Genevieve Lakier, and Chenhao Tan. "Community Guidelines Make this the Best Party on the Internet": An In-Depth Study of Online Platforms' Content Moderation Policies. CHI 2024.

C. Estelle Smith, Shamika Klassen, Gale H. Prinster, Chenhao Tan, and Brian C. Keegan. Governance of the Black Experience on Reddit: r/BlackPeopleTwitter as a Case Study in Supporting Sense of Virtual Community for Black Users. CSCW 2024 (**DEI Recognition Award**, Top 4% of Papers).

Gale H. Prinster, C. Estelle Smith, Chenhao Tan, and Brian C. Keegan. *Community Archetypes: An Empirical Framework for Guiding Research Methodologies to Reflect User Experiences of Sense of Virtual Community*. CSCW 2024.

Dang Nguyen, Chacha Chen, He He, and Chenhao Tan. *Pragmatic Radiology Report Generation*. ML4H 2023.

Xiaoming Shi, Siqiao Xue, Kangrui Wang, Fan Zhou, James Y. Zhang, Jun Zhou, Chenhao Tan, and Hongyuan Mei. *Language Models Can Improve Event Prediction by Few-Shot Abductive Reasoning*. NeurIPS 2023.

Karen Zhou and Chenhao Tan. *Entity-Based Evaluation of Political Bias in Automatic Summarization*. Findings of EMNLP 2023 (Short Paper). Presented at TrustNLP at ACL 2023.

Nan-Jiang Jiang, Chenhao Tan, and Marie-Catherine de Marneffe. *Ecologically Valid Explanations for Label Variation in NLI*. Findings of EMNLP 2023 (Short Paper).

Vivian Lai, Yiming Zhang, Chacha Chen, Q. Vera Liao, and Chenhao Tan. *Selective Explanations: Leveraging Human Input to Align Explainable AI*. CSCW 2023.

Raul Castro Fernandez, Aaron J. Elmore, Michael J. Franklin, Sanjay Krishnan, and Chenhao Tan. *How large language models will disrupt data management*. VLDB 2023.

Chacha Chen, Shi Feng, Amit Sharma, and Chenhao Tan. *Machine Explanations and Human Understanding*. Transactions on Machine Learning Research (TMLR); FAccT 2023; **Best Paper Award** at the ICML 2022 Workshop on Human-Machine Collaboration and Teaming.

Vivian Lai, Chacha Chen, Alison Smith-Renner, Q. Vera Liao, and Chenhao Tan. *Towards a Science of Human-AI Decision Making: A Survey of Empirical Studies*. FAccT 2023.

Mourad Heddaya, Solomon Dworkin, Chenhao Tan, Rob Voigt, and Alexander Zentefis. *Language of Bargaining*. ACL 2023.

Yangqiaoyu Zhou, Yiming Zhang, and Chenhao Tan. *FLamE: Few-shot Learning from Natural Language Explanations*. ACL 2023.

Han Liu, Yizhou Tian, Chacha Chen, Shi Feng, Yuxin Chen, and Chenhao Tan. *Learning Human-Compatible Representations for Case-Based Decision Support*. In Proceedings of ICLR 2023. Presented at the ICML 2022 Workshop on Human-Machine Collaboration and Teaming.

Yiming Zhang, Yangqiaoyu Zhou, Samuel Carton, and Chenhao Tan. *Learning to Ignore Adversarial Attacks*. In Proceedings of EACL 2023.

Yuan Yuan, Tracy Xiao Liu, Chenhao Tan, Qian Chen, Alex ‘Sandy’ Pentland, and Jie Tang. *Gift Contagion in Online Groups: Evidence From Virtual Red Packets*. Management Science. 2023.

Xuan Yang, Yang Yang, Chenhao Tan, Yinghe Lin, Zhengzhe Fu, Fei Wu, and Yueting Zhuang. *Unfolding and Modeling the Recovery Process after COVID Lockdowns*. Scientific Reports. 2023.

Abhinav Kumar, Chenhao Tan, and Amit Sharma. *Probing Classifiers are Unreliable for Concept Removal and Detection*. NeurIPS 2022.

Yiming Zhang, Shi Feng, and Chenhao Tan. *Active Example Selection for In-Context Learning*. EMNLP 2022.

Vivian Lai, Samuel Carton, Rajat Bhatnagar, Q. Vera Liao, Yunfeng Zhang, and Chenhao Tan. *Human-AI Collaboration via Conditional Delegation: A Case Study of Content Moderation*. In Proceedings of the ACM CHI Conference on Human Factors in Computing Systems 2022 (CHI 2022).

Samuel Carton, Surya Kanoria, and Chenhao Tan. *What to Learn, and How: Toward Effective Learning from Rationales*. In Findings of ACL 2022.

Chenhao Tan. *On the Diversity and Limits of Human Explanations*. NAACL 2022.

Jamar L. Sullivan, Will Brackenbury, Andrew McNutt, Kevin Bryson, Kwam Byll, Yuxin Chen, Michael Littman, Chenhao Tan, and Blase Ur. *Explaining Why: How Instructions and User Interfaces Impact Annotator Rationales When Labeling Text Data*. In Proceedings of the 2022 Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL 2022).

C. Estelle Smith, Irfanul Alam, Chenhao Tan, Brian C. Keegan, and Anita Blanchard. *The Impact of Governance Bots on Sense of Virtual Community: Development and Validation of the GOV-BOTs Scale*. In Proceedings of the 25th ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW 2022).

Chao-Chun Hsu and Chenhao Tan. *Decision-focused Summarization*. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP 2021).

Han Liu, Vivian Lai, and Chenhao Tan. *Understanding the Effect of Out-of-distribution Examples and Interactive Explanations on Human-AI Decision Making*. In Proceedings of the 2021 Conference on Computer-Supported Cooperative Work and Social Computing (CSCW 2021).

Madhusudhan Aithal Mahabhaleshwara and Chenhao Tan. *On Positivity Bias in Negative Reviews*. In Proceedings of the Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021) (Short Paper).

Ramaravind Kommiya Mothilal, Divyat Mahajan, Chenhao Tan, and Amit Sharma. *Towards Unifying Feature Attribution and Counterfactual Explanations: Different Means to the Same End*. In Proceedings of the Fourth AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society (AIES 2021).

Jason Shuo Zhang, Qin Lv, Brian Keegan, and Chenhao Tan. *Understanding the Diverging User Trajectories in Highly-related Online Communities during the COVID-19 Pandemic*. In Proceedings of the 15th International AAAI Conference on Web and Social Media (ICWSM 2021).

Isil Erel, Lea Stern, Chenhao Tan, and Michael Weisbach. *Selecting Directors Using Machine Learning*. Review of Financial Studies. 2021.

Abhijit Suresh, Jennifer Jacobs, Vivian Lai, Chenhao Tan, Wayne Ward, James H. Martin, and Tamara Sumner. *Using Transformers to Provide Teachers with Personalized Feedback on their Classroom Discourse: The TalkMoves Application*. AAAI Spring Symposium Series 2021.

Chao-Chun Hsu, Shantanu Karnwal, Sendhil Mullainathan, Ziad Obermeyer, and Chenhao Tan. *Characterizing the Value of Information in Medical Notes*. In Findings of EMNLP 2020.

Samuel Carton, Anirudh Rathore, and Chenhao Tan. *Evaluating and Characterizing Human Rationales*. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP 2020).

Vivian Lai, Han Liu, and Chenhao Tan. *“Why is ‘Chicago’ deceptive?”: Towards Building Model-Driven Tutorials for Humans*. In Proceedings of the ACM CHI Conference on Human Factors in Computing Systems 2020 (CHI 2020).

Ramaravind Kommiya Mothilal, Amit Sharma, and Chenhao Tan. *Explaining Machine Learning Classifiers through Diverse Counterfactual Explanations*. In Proceedings of ACM FAccT Conference 2020.

Vivian Lai, Samuel Carton, and Chenhao Tan. *Harnessing Explanations to Bridge AI and Humans*. In Fair & Responsible AI Workshop, CHI 2020.

Brian Lubars and Chenhao Tan. *Ask not what AI can do for you, but what AI should do: Towards a framework of task delegability*. In Proceedings of the Thirty-Third Conference on Neural Information Processing Systems (NeurIPS 2019) (Spotlight).

David Atkinson, Kumar Bhargav Srinivasan, and Chenhao Tan. *What Gets Echoed? Understanding the ‘Pointers’ in Explanations of Persuasive Arguments*. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP 2019).

Vivian Lai, Jon Z. Cai, and Chenhao Tan. *Many Faces of Feature Importance: Comparing Built-in and Post-hoc Feature Importance in Text Classification*. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP 2019).

Vivian Lai and Chenhao Tan. *On Human Predictions with Explanations and Predictions of Machine Learning Models: A Case Study on Deception Detection*. In Proceedings of ACM FAccT Conference 2019.

Kelvin Luu, Chenhao Tan, and Noah A. Smith. *Measuring Online Debaters’ Persuasive Skill from Text over Time*. Transactions of the Association for Computational Linguistics (TACL 2019).

Xiaochuang Han, Eunsol Choi, and Chenhao Tan. *No Permanent Friends or Enemies: Tracking Relationships between Nations from News*. In Proceedings of the 2019 Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL 2019).

Jason Shuo Zhang, Chenhao Tan, and Qin Lv. *Intergroup Contact in the Wild: Characterizing Language Differences between Intergroup and Single-group Members in NBA-related Discussion Forums*. In Proceedings of the Conference on Computer-Supported Cooperative Work and Social Computing (CSCW 2019).

Kumar Bhargav Srinivasan, Cristian Danescu-Niculescu-Mizil, Lillian Lee, and Chenhao Tan. *Content Removal as a Moderation Strategy: Compliance and Other Outcomes in the ChangeMyView Community*. In Proceedings of the Conference on Computer-Supported Cooperative Work and Social Computing (CSCW 2019).

Tiago Cunha, David Jurgens, Chenhao Tan, and Daniel Romero. *Are All Successful Communities Alike? Characterizing and Predicting the Success of Online Communities*. In Proceedings of The Web Conference 2019 (WWW 2019).

Ziqiang Cheng, Yang Yang, Chenhao Tan, Denny Cheng, Alex Cheng, and Yueting Zhuang. *What Makes a Good Team? A Large-scale Study on the Effect of Team Composition in Online Gaming*. In Proceedings of The Web Conference 2019 (WWW 2019) (Short Paper).

Jason Shuo Zhang, Chenhao Tan, and Qin Lv. *"This is why we play": Characterizing Online Fan Communities of the NBA Teams*. In Proceedings of the Conference on Computer-Supported Cooperative Work and Social Computing (CSCW 2018).

Tal August, Nigini Oliveira, Chenhao Tan, Noah A. Smith, and Katharina Reinecke. *Framing Effect: The Choice of Slogans Used to Advertise Online Experiments Can Boost Recruitment and Lead to Sample Biases*. In Proceedings of the Conference on Computer-Supported Cooperative Work and Social Computing (CSCW 2018).

Dallas Card, Chenhao Tan, and Noah A. Smith. *Neural Models for Documents with Metadata*. In Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL 2018).

Nelson F. Liu, Omer Levy, Roy Schwartz, Chenhao Tan, and Noah A. Smith. *LSTMs Exploit Linguistic Attributes of Data*. In Proceedings of ACL 2018 RepL4NLP Workshop (**Best Paper Award**).

Chenhao Tan. *Tracing Community Genealogy: How New Communities Emerge from the Old*. In Proceedings of the 12th International AAAI Conference on Web and Social Media (ICWSM 2018).

Chenhao Tan, Hao Peng, and Noah A. Smith. *"You are no Jack Kennedy": On Media Selection of Highlights from Presidential Debates*. In Proceedings of The Web Conference (WWW 2018).

Yang Yang, Zongtao Liu, Chenhao Tan, Fei Wu, Yueting Zhuang, and Yafeng Li. *To Stay or to Leave: Churn Prediction for Urban Migrants in the Initial Period*. In Proceedings of The Web Conference (WWW 2018).

Isil Erel, Léa H. Stern, Chenhao Tan, and Michael S. Weisbach. *Could Machine Learning Help Companies Select Better Board Directors?* Harvard Business Review. April, 2018.

Elizabeth Clark, Anne Ross, Chenhao Tan, Yangfeng Ji, and Noah A. Smith. *Creative Writing with a Machine in the Loop: Case Studies on Slogans and Stories*. In Proceedings of the 23rd ACM Conference on Intelligent User Interfaces (IUI 2018).

Yang Yang, Chenhao Tan, Zongtao Liu, Fei Wu, and Yueting Zhuang. *Urban Dreams of Migrants: A Case Study of Migrant Integration in Shanghai*. In Proceedings of the 32nd AAAI Conference on Artificial Intelligence (AAAI 2018).

Yangfeng Ji, Chenhao Tan, Sebastian Martschat, Yejin Choi, and Noah A. Smith. *Dynamic Entity Representations in Neural Language Models*. In Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing (EMNLP 2017).

Chenhao Tan, Dallas Card, and Noah A. Smith. *Friendships, Rivalries, and Trysts: Characterizing Relations between Ideas in Texts*. In Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (ACL 2017).

Chenhao Tan, Vlad Niculae, Cristian Danescu-Niculescu-Mizil, and Lillian Lee. *Winning arguments: Interaction dynamics and persuasion strategies in good-faith online discussions*. In Proceedings of the 25th International World Wide Web Conference (WWW 2016).

Isabel Kloumann, Chenhao Tan, Jon Kleinberg, and Lillian Lee. *Internet Collaboration on Extremely Difficult Problems: Research versus Olympiad Questions on the Polymath Site*. In Proceedings of the 25th International World Wide Web Conference (WWW 2016).

Chenhao Tan, Adrien Friggeri, and Lada Adamic. *Lost in propagation? Unfolding news cycles from the source*. In Proceedings of the 10th International AAAI Conference on Weblogs and Social Media (ICWSM 2016).

Jack Hessel, Chenhao Tan, and Lillian Lee. *Science, AskScience, and BadScience: On the Coexistence of Highly Related Communities*. In Proceedings of the 10th International AAAI Conference on Weblogs and Social Media (ICWSM 2016).

Chenhao Tan and Lillian Lee. *All Who Wander: On the Prevalence and Characteristics of Multi-community Engagement*. In Proceedings of the 24th International World Wide Web Conference (WWW 2015).

Chenhao Tan, Lillian Lee, and Bo Pang. *The effect of wording on message propagation: Topic- and author-controlled natural experiments on Twitter*. In Proceedings of the Association for Computational Linguistics (ACL 2014).

Chenhao Tan and Lillian Lee. *A Corpus of Sentence-level Revisions in Academic Writing: A Step towards Understanding Statement Strength in Communication*. In Proceedings of the Association for Computational Linguistics (ACL 2014) (Short Paper).

Jon Kleinberg, Sendhil Mullainathan, Chenhao Tan, and Thomas Zimmermann. *Inductive Testing: Theory and an Application to the Disposition Effect*. NBER Behavioral Finance Working Group, 2014.

Chenhao Tan, Ed H. Chi, David Huffaker, Gueorgi Kossinets, and Alexander J. Smola. *Instant Foodie: Predicting Expert Ratings From Grassroots*. In Proceedings of the 22nd ACM International Conference on Information and Knowledge Management (CIKM 2013).

Daniel M. Romero, Chenhao Tan, and Johan Ugander. *On the Interplay between Social and Topical Structure*. In Proceedings of the 7th International AAAI Conference on Weblogs and Social Media (ICWSM 2013).

Eunsol Choi, Chenhao Tan, Lillian Lee, Cristian Danescu-Niculescu-Mizil, and Jennifer Spindel. *Hedge detection as a lens on framing in the GMO debates: A position paper*. In Proceedings of the Workshop on Extra-Propositional Aspects of Meaning in Computational Linguistics (2012).

Chenhao Tan, Evgeniy Gabrilovich, and Bo Pang. *To Each His Own: Personalized Content Selection based on Text Comprehensibility*. In Proceedings of the Fifth ACM International Conference on Web Search and Data Mining (WSDM 2012).

Wenbin Tang, Jie Tang, Tao Lei, Chenhao Tan, Bo Gao, and Tian Li. *On optimization of expertise matching with various constraints*. In Neurocomputing (July, 2012).

Chenhao Tan, Lillian Lee, Jie Tang, Long Jiang, Ming Zhou, and Ping Li. *User-level sentiment analysis incorporating social networks*. In Proceedings of the Seventeenth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (SIGKDD 2011).

Bin Lu, Chenhao Tan, Claire Cardie, and Benjamin K. Tsou. *Joint Bilingual Sentiment Classification with Unlabeled Parallel Corpora*. In Proceedings of the Association for Computational Linguistics (ACL 2011).

Shaomei Wu, *Chenhao Tan*, Jon Kleinberg, and Michael Macy. *Does Bad News Go Away Faster?* In Proceedings of the Fifth International AAAI Conference on Weblogs and Social Media (ICWSM 2011) (Short Paper).

Fei Wang, *Chenhao Tan*, Arnd Christian König, and Ping Li. *Efficient Document Clustering via Online Nonnegative Matrix Factorizations*. In Proceedings of the SIAM International Conference on Data Mining (SDM 2011).

Chenhao Tan, Jie Tang, Jimeng Sun, Quan Lin, and Fengjiao Wang. *Social Action Tracking via Noise Tolerant Time-varying Factor Graphs*. In Proceedings of the Sixteenth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (SIGKDD 2010).

Wenbin Tang, Jie Tang, and *Chenhao Tan*. *Expertise Matching via Constrained-based Optimization*. In Proceedings of the 2010 IEEE/WIC/ACM International Conferences (WIC 2010).

Invited Talks

Three Acts of AI Scientists.

MIT, 2026.

University of Chicago, 2026.

NICE, 2026.

Northwestern University (Ryan Institute on Complexity), 2026.

Science in the Age of AI.

Allen Institute for Artificial Intelligence (AI2), 2026.

Northwestern University (Kellogg), 2026.

University of Chicago, 2026.

First Principles, 2026.

DLCT, 2026.

Keynote at the Tenth International Workshop on Symbolic-Neural Learning (SNL), 2026.

SkAI, 2025.

ICAIS, 2025.

Alignment Beyond Human Preferences: Use Human Goals to Guide AI towards Complementary AI.

University of Colorado Boulder, 2025.

Johns Hopkins University, 2024.

Towards Human-Centered AI: How to Generate Useful Explanations for Human-AI Decision Making.

Carnegie Mellon University, 2024.

Northwestern University (Technology and Social Behavior Colloquium), 2024.

University of Chicago, 2024.

ETH Zurich, 2024.

Columbia University, 2023.

Georgia Tech, 2023.

MIT, 2023.

University of Washington, 2023.

Towards Human-Centered AI: Understanding and Improving Human Decision Making through Explanations.

Northwestern University (Kellogg), 2024.

Machine Learning in Economics Summer Institute, 2023.

Three Lessons Towards Human-Centered Explanations.

Keynote at the Artificial Intelligence & Human-Computer Interaction Workshop at ICML, 2023.

The Future of AI is Human-Centered.

Keynote at the Research Day at Governors State University, 2023.

Towards Human-Centered Explanations of AI Predictions.

Stanford University, 2023.

University of Cambridge, 2023.

University of Illinois Chicago, 2023.

University of Pennsylvania, 2022.

University of Toronto, 2022.

New York University, 2022.

University of Texas at Austin, 2022.

TTIC, 2022.

Tutorial on Human-Centered Evaluation of Explanations.

NAACL, 2022.

Language, Information, and Decision Making.

Machine Learning in Economics Summer Institute, 2022.

Brown Bag at Booth Healthcare Initiative, 2022.

Why are Explanations Falling Short of their Promises?

AAAI Workshop on Explainable Agency in Artificial Intelligence, 2022.

Towards Human-Centered AI: Complementing Humans with AI.

The Academic Fringe Festival on Designing at Scale with Human-AI Collaboration, 2022.

Characterizing the “Value” of Information.

TTIC, 2021.

Towards Human-Centered AI: Harnessing Explanations to Improve Human Decision Making in Discovery Tasks.

The Second ELLIS HCML Workshop at ICLR, 2021.

Towards Human-Centered AI: Harnessing Explanations to Improve Human Decision Making.

University of Wisconsin-Madison, 2021.

- Lunch for CAPP at the University of Chicago, 2021.
- A Multitude of Communities: Genealogy and Intergroup Contact.
University of Notre Dame, 2021.
- Towards Human-Centered AI: Understanding Human Production and Consumption of Explanations.
University of Maryland, 2021.
University of Michigan, 2021.
- Towards Human-Centered AI: Understanding Language in Social Contexts.
University of Chicago, 2020.
Carnegie Mellon University, 2020.
Google, 2020.
Microsoft Research, 2020.
Yale University, 2020.
University of Illinois Urbana-Champaign, 2020.
University of Washington, 2020.
ASAPP, 2019.
- Towards Human-Centered AI: Harness AI to Improve Human Decision Making.
Harris School of Public Policy, University of Chicago, 2020.
Tsinghua University, 2020.
USC AI Rising Stars Symposium, 2019.
- Towards Human-Centered AI: Opportunities and Challenges for Natural Language Processing in Policy.
Deploying Machine Learning Tools for Public Policy Impact, APPAM Pre-Conference Workshop, 2019.
- The Promise of Human-Centered AI for Addressing Challenges in Literacy.
Facebook, 2019.
- The Future of Brand Engagement: Community, Trust and Digital Discourse.
Reddit, 2018.
- Towards Understanding and Generating Language for Effective Social Interaction.
IST Seminar, Caltech, 2018.
- Characterizing the Ecosystem of Ideas in Texts.
Summer Institute in Computational Social Science, University of Colorado Boulder, 2018.
The Fall 2018 Workshop on Quantitative Methods, University of Colorado Boulder, 2018.
Department of Computer Science, Tsinghua University, 2018.
Computational Semantics Group, University of Colorado Boulder, 2017.
UW-MS NLP Symposium, 2017.

Winning Arguments: Interaction Dynamics and Persuasion Strategies in Good-Faith Online Discussions.

The 24th Annual Northwest Dispute Resolution Conference, 2018.

Colorado Data Science Team, University of Colorado Boulder, 2017.

Phrasing Matters: An Interactive Session on Message Virality Based on Twitter.

The 24th Annual Northwest Dispute Resolution Conference, 2018.

Wandering Individuals and Emerging Communities in a World of Communities.

Department of Sociology, Tsinghua University, 2018.

INFO/HCC Seminar, University of Colorado Boulder, 2017.

It Takes Two to Tango: Understanding the Effects of Wording through Online Interactions.

Fudan University, 2018.

Toutiao AI Lab, 2018.

Zhejiang University, 2017.

A Short Talk on Research.

Zhejiang University, 2017.

Online Social Interactions: A Lens on Humans and a World for Humans.

CSSS Seminar, University of Washington, 2016.

MIT EECS, 2016.

CMU Computer Science Department, 2016.

CMU Heinz College, 2016.

Max Planck Institute for Software Systems, 2016.

University of Michigan, 2016.

University of Colorado Boulder, 2016.

Rice University, 2016.

Northeastern University, 2016.

It Takes Two to Tango: Understanding the Effects of Language via “Natural Experiments”.

AI Seminar, Cornell University, 2014.

The Effect of Wording on Message Propagation: Topic- and Author-Controlled Natural Experiments on Twitter.

Tsinghua University, 2014.

More than Accuracy: Interpretability.

Machine Learning Reading Group, Cornell University, 2013.

User-Level Sentiment Analysis Incorporating Social Networks.

Tsinghua University, 2011.

Microsoft Research Asia NLP Group Seminar, 2011.

Professional Service

Co-organizer of the AI and Social Science Conference at the University of Chicago (2026)

Steering Committee of FAccT (2023-2026)

Faculty Advisory Board for Cognitive Science, University of Chicago (09/2021-present)

Advisory Board for the Cureus Journal of AI-Augmented Research (2025-present)

Executive Committee Member at the Institute of Cognitive Science, University of Colorado Boulder (10/2019-12/2020)

General Chair:

2023: FAccT

Program Chair:

2022: IC2S2

Senior Area Chair:

2027: NAACL

2026: ACL

2025: ACL

2024: EMNLP

2023: ACL

2022: AACL

2021: NAACL

Area Chair (Senior PC):

2026: COLM

2025: COLM

2021: ACL, CSCW, AKBC

2020: EMNLP

2019: NAACL, EMNLP, AAAI

Tutorial Chair: ICWSM 2020

Workshop Organization:

Interp4Discovery at NeurIPS 2026

NLP4Democracy at COLM 2025

The First Workshop on AI Behavioral Science (AIBS 2024)

Beyond Online Data Workshop at ICWSM 2018

PC Member:

2026: ICML, ARR

2025: ARR

2022: NAACL, FAccT, EMNLP

2021: CSCW

2020: WWW

2019: WWW, FAT*, ACL, NeurIPS (one of 400 highest-scoring reviewers)

2018: WWW, NAACL, ACL, KDD, NeurIPS, ICWSM, BigNet Workshop

2017: WWW, WSDM, KDD, ACL, EMNLP, CoNLL

2016: WWW, KDD, WSDM, IJCAI, NAACL, ICWSM, SDM, CIKM

2015: WWW, KDD, WSDM

2014: WWW, EMNLP, ICWSM, DYAD

2013: IJCAI, IJCNLP, NAACL Student Research Workshop

2012: ADMA

2011: ADMA

Panelist: NSF (2025, 2024, 2023, 2021, 2020, 2019, 2019)

Session Chair: NAACL (2019), WWW (2015)

Reviewer: Nature Communications (2025), JMLR (2023, 2021), PNAS (2020), ICWSM (2019), HICSS (2019), TACL (2020, 2019, 2018, 2017, 2016), Transactions on Knowledge Discovery from Data (2015, 2012), TOIS (2018), PLOS ONE (2018), CHI (2017), CSCW (2016), WWW (2011, 2010), ICDM (2011), KDD (2010), CIKM (2010)

Selected Media Coverage

Science (2026): Researchers presented 6000 papers at a major meeting. Could AI reproduce their findings?

Nature (2026): No humans allowed: scientific AI agents get their own social network

UChicago News (2025): Why can't powerful AIs learn basic multiplication?

UChicago News (2024): UChicago Researchers Demonstrate the Quantifiable Uniqueness of Former President Donald Trump's Language Use

The New York Times (2023): How Easy Is It to Fool A.I.-Detection Tools?

PolitiFact (2023): How improperly using AI could deter you from voting, or hurt your health

CU Boulder Today (2020): As the coronavirus spread, 2 social media communities drifted apart

The New York Times (2019): How to Argue on the Internet Without Losing Your Mind

9NEWS (2019): Study of NBA Reddit forum hopes to learn more about how people interact with each other online

CU Boulder Today (2019): Flagrant fouls: What Reddit's basketball fans can tell us about online discourse

MIT Technology Review (2019): WeChat is running a natural experiment in human generosity

Harvard Business Review (2018): Want Less-Biased Decisions? Use Algorithms.

NPR (2017): Change My View On Reddit Helps People Challenge Their Own Opinions

New Scientist (2017): Phone metadata reveals where city migrants go and who they call
You Are Not So Smart podcast (2016): Change My View
Slate (2016): You Can't Win: How algorithms are changing the ways we argue
The Washington Post (2016): How to change someone's mind, according to science
New York Magazine (2016): A Subreddit Sparked a Scientific Inquiry Into How to Change Someone's Mind
Daily Mail (2016): How to win ANY argument using science: Experts reveal which words to use and how one simple trick can help you get your own way
The New York Times (2014): A 25-Question Twitter Quiz to Predict Retweets
The Washington Post (2014): This is the perfect tweet, according to science
The Physics arXiv Blog (2014): Data Mining Reveals How Wording Influences Tweet Propagation
Engadget (2014): Crafting the perfect tweet is one-third talent, two-thirds science
The Daily Dot (2014): A data miner's guide to crafting the perfect tweet

Postdoctorates

Davi Bastos Costa
Xiao Liu
Cristina Garbacea (2026) → CISA Helmholtz Center for Information Security
Julia Mendelsohn (2025) → Assistant Professor at University of Maryland
Shi Feng (2023) → Research Associate at New York University → Assistant Professor at George Washington University
Sam Carton (2022) → Assistant Professor at the University of New Hampshire
Estelle Smith (2022) → Assistant Professor at Colorado School of Mines

PhD Students

Xiaoyan Bai
Chacha Chen (2025) → Apple → MAI
Qirun Dai
Mourad Heddaya (2026)
Chao-Chun Hsu (2024) → Amazon
Han Liu (2024) → EvenUp
Haokun Liu
Dang Nguyen

Karen Zhou (2026) → Google Research

Yangqiaoyu (Rosa) Zhou (2025) → Microsoft

Vivian Lai (2022) → Visa Research → Microsoft

Jason Shuo Zhang (2020, co-advised with Qin Lv) → postdoc at the Max Planck Institute → ASOS → Tripadvisor

Master Students (thesis/independent study)

Yi Wu (class of 2026, University of Chicago)

Yuyang Jiang (class of 2025, University of Chicago)

Tejes Srivastava (class of 2025, University of Chicago)

Yixuan Wang (class of 2025, University of Chicago)

Jessica Ji (class of 2024, University of Chicago)

Jiamin Yang (class of 2024, University of Chicago)

Zander Meitus (class of 2023, University of Chicago)

Harry Tian (class of 2023, University of Chicago)

Yiming Zhang (class of 2023, University of Chicago)

Madhusudhan Aithal Mahabhaleshwara (class of 2021, University of Colorado Boulder)

Rajat Bhatnagar (class of 2021, University of Colorado Boulder)

Anirudh Rathore (class of 2021, University of Colorado Boulder)

Surya Kanoria (class of 2020, University of Colorado Boulder)

Brian Lubars (class of 2019, University of Colorado Boulder)

Shantanu Karnwal (class of 2019, University of Colorado Boulder)

Hunter Wapman (class of 2019, University of Colorado Boulder)

Kumar Srinivasan (class of 2019, University of Colorado Boulder)

Vivian Lai (class of 2018, University of Colorado Boulder)

Undergraduate/High School Advising

Lawrence Lu (class of 2028, University of Chicago)

Angus Chen (class of 2028, University of Chicago)

Christian Lee (class of 2027, University of Chicago)

Mingxuan Li (class of 2027, University of Chicago)

Justin Hwang (class of 2027, University of Chicago)

Julia Margie (class of 2026, University of Chicago)
Aditya Singh (class of 2026, University of Chicago)
Sohini Banerjee (class of 2026, University of Chicago)
Miles Wang (class of 2026, Harvard University)
Tianyu Cao (class of 2024, Zhejiang University)
Lily Wang (class of 2024, University of Chicago)
Dang Nguyen (class of 2024, University of Chicago)
Grace Wang (class of 2024, University of Chicago)
Aditya Krishna (class of 2024, University of Chicago)
Milo Chase (class of 2023, University of Chicago)
Kenan Tang (class of 2023, University of Chicago)
Anne Mykland (class of 2023, University of Chicago Laboratory School)
Brent DeVries (class of 2022, University of Chicago)
Ava Bartolome (class of 2022, Clark University → PhD at the University of Minnesota)
Qitian Hu (class of 2022, University of Chicago)
Martin Chou (class of 2021, Tsinghua University)
Yangqiaoyu Zhou (class of 2020, Carleton College)
Jie Wang (class of 2020, University of Colorado Boulder)
Zhenqi Li (class of 2020, University of Colorado Boulder)
Xiaochuang Han (class of 2019, Georgia Tech)
David Atkinson (class of 2019, University of Colorado Boulder)
Cora Schneck (class of 2018, University of Colorado Boulder)
Kelvin Luu (class of 2015, Cornell University)
Eunsol Choi (class of 2012, Cornell University, honorable mention for the CRA Outstanding Undergraduate Research Award in 2011)

Committee Members

I served as a committee member for the following students: Harvey Fu (PhD, University of Chicago), Todd Nief (PhD, University of Chicago), Zain Sarwar (PhD, University of Chicago), Vishakh Padmakumar (PhD, New York University), Jiaming Qu (PhD, University of North Carolina), Afia Khan (MD/PhD, University of Chicago), Nan-Jiang Jiang (PhD, Ohio State University), Joon Sik Kim (PhD, Carnegie Mellon University), Valdemar Danry (master, MIT), Gil Rosenthal (master, University of Chicago), Jingwen Jiang (PhD, University of Chicago), Divyansh Kaushik (PhD, Carnegie Mellon University), Rehan Ahmed (PhD, University of Colorado Boulder), Anupama (master, University of Colorado Boulder), Jon Cai (PhD, University of Colorado Boulder), Nora Connor (PhD, University of Colorado Boulder), Xiaolan Ding (PhD, University of Chicago), Xu Han (PhD, University of Colorado Boulder), Shudong Hao (PhD, University of Colorado Boulder), Daniel Leopold (Psychology PhD, University of Colorado Boulder), Forough Poursabzi (PhD, University of Colorado Boulder), Tyler Scott (undergraduate, University of Colorado Boulder), Saumya Sinha (master, University of Colorado Boulder), Hayeong Song (master, University of Colorado Boulder), Angela Stewart (PhD, University of Colorado Boulder), Shubha Swamy (undergraduate, University of Colorado Boulder), Yichen Wang (PhD, University of Colorado Boulder), Hunter Wapman (PhD, University of Colorado Boulder), Max Wenzel (undergraduate, University of Colorado Boulder), Mark Wilmes (undergraduate, University of Colorado Boulder), and Linzi Xing (master, University of Colorado Boulder).

Industrial Experience

Visiting Scientist, 2024-2025

Abridge, Remote

Data Scientist Intern, 06/2015-08/2015

Facebook, Menlo Park, CA

Understanding news propagation.

Research Intern, 05/2012-08/2012

Google, Mountain View, CA

Harmonizing place scores and instant foodie.

Research Intern, 06/2011-08/2011

Yahoo! Research, Santa Clara, CA

Personalization based on text comprehensibility.

Research Intern, 04/2010-07/2010

Microsoft Research Asia, Beijing, China

Sentiment analysis incorporating social networks.

Teaching Experience

Instructor, Winter 2021, Winter 2022, Winter 2026

University of Chicago, Chicago, IL

Natural Language Processing.

Instructor, Spring 2026

University of Chicago, Chicago, IL

Large Language Models.

Instructor, Spring 2022, Spring 2023, Spring 2024

University of Chicago, Chicago, IL
Machine Learning for Public Policy.

Instructor, Spring 2021, Spring 2022, Spring 2023, Spring 2024
University of Chicago, Chicago, IL
Human-Centered Machine Learning.

Instructor, Fall 2018, Fall 2020
University of Colorado, Boulder, CO
Undergraduate Machine Learning, CSCI 4622.

Instructor, Spring 2018, Spring 2019
University of Colorado, Boulder, CO
Human-Centered Machine Learning, CSCI 7000.

Instructor, Fall 2017, Fall 2019
University of Colorado, Boulder, CO
Machine Learning, CSCI 5622.

Teaching Assistant, Fall 2013
Cornell University, Ithaca, NY
Information Retrieval, CS4300/INFO4300.

Teaching Assistant, Fall 2011
Cornell University, Ithaca, NY
Machine Learning, CS4780/5780.